

Multi-Threaded Web Server with IPC and Semaphores

SO 25/26

Design Document


Department of Electronics, Telecommunications and Informatics

Bernardo Mota Coelho - 125059
Tiago Francisco Crespo do Vale - 125913

9/12/2025



1. Architecture diagram




2. Flowcharts

Flowcharts


This document includes the following specific flowcharts:

3. Cache Flowchart
4. Config Flowchart
5. HTTP Flowchart
6. Main Flowchart
7. Master Flowchart
8. Semaphores Flowchart
9. Shared Memory Flowchart
10. Stats Flowchart
11. ThreadPool Flowchart
12. Worker Flowchart


Cache FlowChart:




Config FlowChart:




HTTP FlowChart:




Main FlowChart:




Semaphores FlowChart:




Master FlowChart:




Shared Memory FlowChart:




Stats FlowChart:



ThreadPool FlowChart:



Worker FlowChart:



3. Synchronization Analysis

This HTTP server uses multiple synchronization mechanisms across processes and threads to ensure data consistency and avoid race conditions.

Process-Level Syncronization (IPC)

Shared Memory (shm_data)

```
// Master creates, workers attach
shared_data_t* shm_data = shm_create(); // master.c
shared_data_t* shm_data = shm_create(); // worker.c (attaches)
```

Data shared:

server_stats_t stats	- Request Statistics
shared_queue_t queue	- Work queue (if used)

Protection: POSIX named semaphores

Thread-Level Synchronization

A. Thread Pool Queue (thread_pool.c)

```

pthread_mutex_t mutex;
pthread_cond_t cond_non_empty; // Consumers wait
pthread_cond_t cond_non_full; // Producers wait

// Producer: thread_pool_add()
pthread_mutex_lock(&mutex);
while (queue_full) pthread_cond_wait(&cond_non_full, &mutex);
// Add to queue
pthread_cond_signal(&cond_non_empty);
pthread_mutex_unlock(&mutex);

// Consumer: queue_pop()
pthread_mutex_lock(&mutex);
while (!queue_empty) pthread_cond_wait(&cond_non_empty, &mutex);
// Remove from queue
pthread_cond_signal(&cond_non_full);
pthread_mutex_unlock(&mutex);
  
```

B. Cache System (cache.c)

```

pthread_rwlock_t cache_rwlock; // READER-WRITER LOCK

// Readers (cache_get):
pthread_rwlock_rdlock(&cache_rwlock);
// Multiple threads can read simultaneously
pthread_rwlock_unlock(&cache_rwlock);

// Writers (cache_put):
pthread_rwlock_wrlock(&cache_rwlock);
// Only one thread can write, no readers
pthread_rwlock_unlock(&cache_rwlock);
  
```

A: Pattern: Producer-Consumer with bounded buffer

B: Advantage: Better performance for read-heavy workload (web server cache)

Critical Sections Analysis

Section 1: Statistics Updates

```
// stats.c - stats_update()
sem_wait(sem_stats); // ENTER CRITICAL SECTION
stats->total_requests++;
stats->bytes_transferred += bytes;
// ... update status counters
sem_post(sem_stats); // EXIT CRITICAL SECTION
```

Risk: Without semaphore, race condition on counters when multiple threads update simultaneously

Section 2:

```
// logger.c - logger_log()
// Implicit semaphore via sem_ws_log
fprintf(log_fp, "[%s] %s \"%s %s\" %d %ld\n", ...);
fflush(log_fp);
```

Risk: Log entries could interleave without synchronization.

Section 3:

```
// cache.c - cache_put()
pthread_rwlock_wrlock(&cache_rwlock);
if (e->valid && e->data) free(e->data); // Free old
e->data = malloc(size); // Allocate new
memcpy(e->data, data, size); // Copy data
pthread_rwlock_unlock(&cache_rwlock);
```

Risk: Without RW-lock, cache corruption or double-free

Potential Deadlocks

Scenario 1: Nested Locks

```

// Thread 1:
pthread_rwlock_wrlock(&cache_lock); // LOCK A
sem_wait(sem_stats); // LOCK B

// Thread 2:
sem_wait(sem_stats); // LOCK B
pthread_rwlock_wrlock(&cache_lock); // LOCK A

```

Analysis: Potential deadlock if threads acquire locks in different order.

Solution: Always acquire locks in consistent order:

1. Process semaphores first (sem_stats, sem_log)
2. Then thread mutexes/RW-locks

Scenario 2: Conditional Wait Timeout

```

// Thread 1:
pthread_rwlock_wrlock(&cache_lock); // LOCK A
sem_wait(sem_stats); // LOCK B

// Thread 2:
sem_wait(sem_stats); // LOCK B
pthread_rwlock_wrlock(&cache_lock); // LOCK A

```

Prevention: Timeout prevents permanent deadlock if queue stays full.

Race Condition Analysis

Race 1: Cache Read While Writing

```

// Thread A (Writer):
free(old_data); // ← Could free while Thread B reads
malloc(new_size);
memcpy(new_data, ...);

// Thread B (Reader):
char* data = e->data; // ← Could read freed memory
size_t size = e->size;

```

Prevention: RW-lock ensures writers have exclusive access

Race 2: Statistics Counter Increment

```
// Without semaphore:  
Thread A: read total (100)  
Thread B: read total (100)  
Thread A: increment (101)  
Thread B: increment (101) // Lost update!  
Final: 101 (should be 102)
```

Prevention: Semaphore ensures atomic increment.

Race 3: File Serving with Cache

```
// Thread A: Cache miss, reading file  
read(file_fd, file_data, st.st_size); // Reading  
cache_put(path, file_data, size); // Caching  
  
// Thread B: Same request arrives  
cache_get(path, &data, &size); // Might get partial data
```

Prevention: Cache put is atomic under write lock