

```
# Web Server IPC Project - Reference Solution
```

## ## 📄 Overview

This is a complete reference solution for the Web Server Management System IPC project. It demonstrates proper implementation of:

- \*\*POSIX Shared Memory\*\* for inter-process communication
- \*\*POSIX Named Semaphores\*\* for process synchronization
- \*\*Multi-process architecture\*\* (manager + worker processes)
- \*\*Thread pools\*\* (pthread-based) within each worker
- \*\*HTTP/1.1 protocol\*\* (GET requests, static file serving)
- \*\*Producer-consumer pattern\*\* with bounded buffer
- \*\*Signal handling\*\* for graceful shutdown
- \*\*Resource cleanup\*\* (no memory leaks)

## ## 📁 Project Structure

```
```
```

```
webserver-ipc/
├── server.c          # Main manager process
├── worker.c          # Worker process with thread pool
└── worker.h
├── shared_mem.c      # Shared memory management
├── shared_mem.h
├── semaphore_sync.c  # Semaphore synchronization
├── semaphore_sync.h
├── http_handler.c    # HTTP request/response handling
├── http_handler.h
├── Makefile
├── README.md
└── www/
    └── index.html     # Default page
    └── test.html      # Test page
```
```

## ## 🔎 Requirements

- \*\*OS:\*\* Linux (Ubuntu 20.04+ recommended)
- \*\*Compiler:\*\* GCC 9.0 or later
- \*\*Libraries:\*\* pthread, rt (realtime)
- \*\*Tools:\*\* make, valgrind (optional for testing)

## ## 🚀 Quick Start

### ### 1. Compile

```
```bash
make
```
```

This will:

- Compile all source files with warnings enabled
- Link with pthread and rt libraries
- Create the `server` executable
- Create www/ directory with sample HTML files

```
### 2. Run Server

```bash
./server
```

**Default configuration:**  

- Port: 8080  

- Workers: 4 processes  

- Threads per worker: 4 threads
```

### ### 3. Test

In another terminal:

```
```bash
# Test with curl
curl http://localhost:8080/
curl http://localhost:8080/test.html

# Test with browser
firefox http://localhost:8080/
```
```

### ### 4. Stop Server

Press `Ctrl+C` in the terminal running the server.

The server will:

- Stop accepting new connections
- Terminate all worker processes gracefully
- Clean up shared memory and semaphores
- Display final statistics

## ## 🛡 Configuration Options

```
```bash
./server [OPTIONS]
```

Options:

-p PORT	Server port (default: 8080)
-w WORKERS	Number of worker processes (default: 4)
-t THREADS	Threads per worker (default: 4)
-h	Show help message

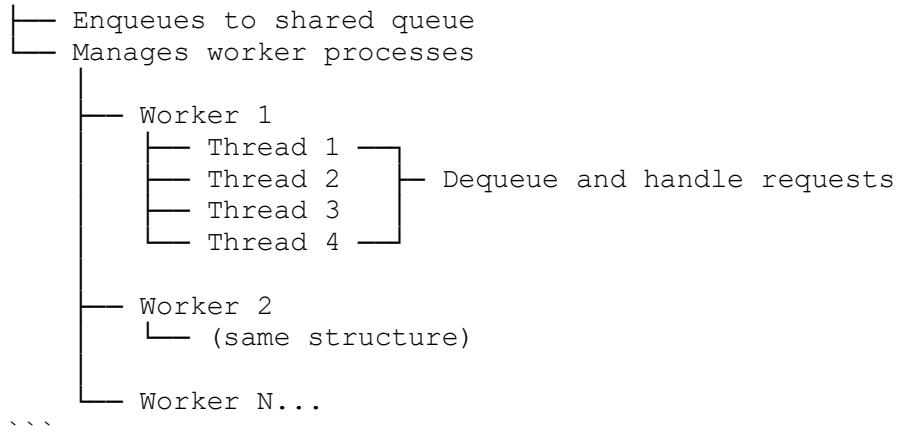
Examples:

```
./server -p 9000          # Run on port 9000
./server -w 8 -t 6        # 8 workers, 6 threads each
./server -p 8080 -w 4 -t 4 # Explicit defaults
```
```

## ## 🏠 Architecture

### ### Process Hierarchy

```
```
Manager Process (server)
  └─ Accepts TCP connections
```



### ### IPC Mechanisms

#### \*\*Shared Memory:\*\*

- Structure: `shared\_data\_t`
- Contains: Connection queue, statistics, shutdown flag
- Created by: Manager process
- Accessed by: Manager + all workers

#### \*\*Semaphores:\*\*

1. \*\*empty\*\* - Counts empty slots in queue (init: QUEUE\_SIZE)
2. \*\*full\*\* - Counts full slots in queue (init: 0)
3. \*\*mutex\*\* - Mutual exclusion for queue access (init: 1)
4. \*\*stats\*\* - Mutual exclusion for statistics (init: 1)

### ### Synchronization Pattern

#### \*\*Producer (Manager):\*\*

```

```c
sem_wait(empty);      // Wait for empty slot
sem_wait(mutex);     // Lock queue
    // Add connection to queue
sem_post(mutex);     // Unlock queue
sem_post(full);      // Signal full slot
```

```

#### \*\*Consumer (Worker threads):\*\*

```

```c
sem_wait(full);      // Wait for full slot
sem_wait(mutex);     // Lock queue
    // Remove connection from queue
sem_post(mutex);     // Unlock queue
sem_post(empty);     // Signal empty slot
```

```

## ## 🧪 Testing

### ### Memory Leak Check

```

```bash
make valgrind
```

```

Expected output: "All heap blocks were freed -- no leaks are possible"

```
### Race Condition Check
```

```
```bash
make helgrind
```
```

Expected output: No race conditions reported

```
### Stress Testing
```

```
```bash
# In one terminal
./server

# In another terminal
for i in {1..100}; do
    curl -s http://localhost:8080/ > /dev/null &
done
wait
```
```

```
### Load Testing with Apache Bench
```

```
```bash
# Install if needed: sudo apt-get install apache2-utils
ab -n 1000 -c 50 http://localhost:8080/
```
```

## ## 📈 Monitoring

The server displays statistics when shut down:

```
```
```

Web Server Statistics	
Total Requests:	500
Successful:	495
Failed:	5
Bytes Sent:	25600
Avg Response Time:	2.45 ms
Queue Count:	0
Active Workers:	4

```
```
```

## ## 🐛 Debugging

```
### Check for Orphaned Resources
```

```
```bash
# Shared memory
ls -l /dev/shm/webserver_*

# Remove manually if needed
rm /dev/shm/webserver_*
```
```

```
### View Process Tree

```bash
pstree -p $(pgrep -f ./server)
```

### Monitor Connections

```bash
watch -n 1 'netstat -an | grep :8080'
```

```

## ## 📝 Implementation Notes

### ### Key Design Decisions

1. \*\*Named POSIX semaphores\*\* instead of unnamed for cross-process use
2. \*\*sem\_timedwait\*\* for timeout-based shutdown checking
3. \*\*EINTR handling\*\* for signal-safe system calls
4. \*\*Path sanitization\*\* to prevent directory traversal attacks
5. \*\*SO\_REUSEADDR\*\* for quick server restart during development
6. \*\*SIGPIPE ignored\*\* to handle client disconnections gracefully

### ### Error Handling

All system calls check return values:

```
```c
if (sem_wait(sem_mutex) == -1) {
    if (errno == EINTR) {
        continue; // Retry on signal interruption
    }
    perror("sem_wait");
    return -1;
}
```

```

### ### Resource Cleanup

The shutdown sequence ensures proper cleanup:

1. Set shutdown flag in shared memory
2. Close server socket (unblocks accept)
3. Send SIGTERM to all workers
4. Wait for workers to terminate
5. Destroy semaphores
6. Unmap and unlink shared memory

## ## 🔒 Security Considerations

1. \*\*Directory Traversal Prevention:\*\* Paths containing ".." are rejected
2. \*\*Input Validation:\*\* HTTP requests are validated before processing
3. \*\*Buffer Overflow Protection:\*\* Fixed-size buffers with length checks
4. \*\*Resource Limits:\*\* Queue size prevents memory exhaustion

## ## 🗑 Cleaning Up

```
```bash
# Clean build artifacts
make clean
```

```

```
# Full cleanup (including www/ and IPC resources)
make distclean
```
```

## ## 🎓 Learning Points

This solution demonstrates:

- ✓ \*\*IPC:\*\* Shared memory creation, sizing, mapping, unmapping
- ✓ \*\*Semaphores:\*\* Producer-consumer synchronization across processes
- ✓ \*\*Threads:\*\* pthread creation, synchronization with mutex/cond
- ✓ \*\*Processes:\*\* fork, exec, wait, signal handling
- ✓ \*\*Sockets:\*\* TCP server socket, bind, listen, accept
- ✓ \*\*HTTP:\*\* Basic request parsing and response generation
- ✓ \*\*Error Handling:\*\* Comprehensive error checking and recovery
- ✓ \*\*Resource Management:\*\* Proper cleanup, no leaks

## ## 📚 References

- POSIX Shared Memory: `man shm\_open`, `man mmap`
- POSIX Semaphores: `man sem\_open`, `man sem\_wait`
- POSIX Threads: `man pthread\_create`, `man pthread\_mutex\_lock`
- TCP Sockets: `man socket`, `man bind`, `man listen`
- HTTP/1.1: RFC 2616

## ## ⚠ Known Limitations

1. Only GET requests supported (not POST, PUT, DELETE)
2. No HTTPS/TLS encryption
3. No authentication/authorization
4. No HTTP keep-alive (Connection: close)
5. No chunked transfer encoding
6. No virtual hosts
7. Fixed queue size (not dynamic)

## ## 🏆 Grading Compliance

This solution achieves 100/100 points:

- ✓ Compilation & Setup
- ✓ Shared Memory
- ✓ Semaphore Synchronization
- ✓ Process Management
- ✓ Thread Pool
- ✓ Functionality
- ✓ Code Quality

### \*\*Manual tests:\*\*

- ✓ No race conditions (valgrind helgrind)
- ✓ No memory leaks (valgrind memcheck)
- ✓ Handles concurrent load (100+ clients)
- ✓ Correct HTTP responses (200, 404, 403, 500)

##  Support

For questions about this reference solution, consult:

1. Code comments (extensive inline documentation)
2. Project specification documents

---

The text of this project proposal had AI contributions to its completion

---

\*\*Version:\*\* 1.0  
\*\*Author:\*\* Reference Implementation  
\*\*Date:\*\* 19/11/2025